

CRSS — Recorded Demo Playbook

Everything you need to record a ~2:30–3:00 demo of CRSS answering Braxton's question — a **video you send** to RA & compliance people who might be interested, drop in DMs, and pin to your profile. Its job isn't to *pitch*; it's to make them **recognise their own problem** so clearly that *they* conclude they need this. You won't be in the room to answer doubts — the video has to do that alone.

Audience: RA officers, compliance leads, consultants, notified-body people — medical devices + AI. Time-poor, skeptical, allergic to hype. **Sincerity beats polish. The tool does the talking; their own recognition does the convincing.**

=====

0. Mindset (read before you record)

- **You are not the regulatory expert — they are.** Your credibility = the tool + your honesty. Lean into *"I built this; you live this; that's why I want your eyes on it."*
 - **No hype, ever.** One overclaim loses this crowd. "Here's what it can't do yet" earns more trust than any superlative.
 - **The 1:40 processing time is a FEATURE for this audience.** An *instant* legal answer is a red flag (that's a model hallucinating). "Slow because it checks itself first" is the sell.
 - **It's sent, not presented — aim for recognition, not persuasion.** They watch alone, probably muted, and can stop any second; you're not there to answer doubts. So you don't *argue* anyone into it — you make them *see their own recurring week* and let "I need this" be their own conclusion. Lead with their problem, prove it works, hand them the mic.
 - **One take, synced voice + screen.** You narrate live over the screen. It's a touch harder (you talk through the wait), but the real-time honesty fits this crowd better than a polished splice — and you can still trim afterward. Rehearse, capture, tidy.
 - **Ship the rough first cut.** The enemy isn't a mediocre video — it's the perfect one you never post. Done > perfect. It gets easier by your third one.
- =====

1. The video at a glance

#	Beat	~Time	What's on screen
1	The hook (the question)	0:00–0:18	The question as text
2	The tension	0:18–0:30	Question, highlight both horns
3	The mirror + what it is	0:30–0:52	"You know this question" → name CRSS + scope, then paste & enter
4	The wait = how it works	0:52–2:15	Draft container (+ optional graph glimpse)
5	The answer (the "aha")	2:15–2:45	Final answer: safety-component + citations, defensibility
6	Honest close + their turn	2:45–3:00	Boundary, then "send me the one that ate <i>your</i> afternoon"

Total ~3:00. The wait-narration (Beat 4) is elastic — it fills whatever the real processing time is, so trim or stretch it to land the whole thing under 3:00 (nearer 2:30 if you can). **Never over 3:00.**

The 3-second rule: on LinkedIn people scroll in 3 seconds. Your first frame + first line must grab. That's why we open on the *question*, not on "Hi, I'm Diego." What CRSS is comes right *after* the hook (Beat 3), wrapped in the **mirror** — "*you know this kind of question*" — so they hear their own week, not a product pitch. Lead with setup and you spend your best 3 seconds on yourself.

2. The full script (voiceover)

Word-for-word "SAY" — adapt to your voice. **Confirm in a dry run that CRSS actually surfaces the "safety component" point and the citations. The real screen is the boss — if it words things differently, match your narration to it.**

Beat 1 — The hook (0:00–0:18)

- **SCREEN:** The question on screen as plain text (don't name who asked it).
- **SAY:** *"A question came up in an EU compliance community, and it looks simple until you try to answer it: a company builds an LLM into a medical device — but for something non-medical, not diagnosis or treatment. Is that LLM a high-risk AI system under the AI Act, or not?"*

Beat 2 — The tension (0:18–0:30)

- **SCREEN:** Keep the question up; highlight the two halves.
- **SAY:** *"You can argue it both ways. The device sits inside the AI Act's Annex I — so maybe high-risk. But the LLM isn't doing anything medical — so maybe not. Which is it?"*

Beat 3 — The mirror, what it is, then ask (0:30–0:52)

- **SCREEN:** Bring CRSS on screen; paste the question, hit enter.
- **SAY:** *"If you work in medtech regulatory or compliance, you know this question — not this exact one, but its shape: the ones that straddle two regulations at once, where every reading is arguable and you lose an afternoon in the AI Act, the MDR, and the guidance just to be sure you're right. That's what I built CRSS for. It answers only from the actual legal text, and it checks its own work — so let me just ask it, and watch what it does."*
- **Why it's shaped this way:** the puzzle (Beats 1–2) stops the scroll; *"you know this question"* makes them supply their **own** example and feel their own week; *"only from the legal text, and checks itself"* is belief-insurance so they trust what they're about to watch. One breath — never a pitch. (Naming the regs also lets the right people self-qualify.)

Beat 4 — The wait = how it works (0:40–2:05) — your best segment

- **SCREEN:** The draft container where you can see it working. Point at it.
- **SAY:** *"Now — it's not instant, and that's the point. While it works, let me show you what's actually happening. Right here, it's writing a first draft — but it won't show you that draft. It takes its own first attempt and does a second pass over it, re-checking every claim against the actual text of the regulations it pulled in. That's deliberate. An instant legal answer is a red flag — that's a model autocompleting from memory, and that's exactly how you get confident, fluent, completely invented citations. This does the opposite: it's slower because it's verifying itself before it shows you a single word. Underneath, it's not treating the regulations as a flat pile of text. [OPTIONAL: cut to graph — see §4] These are the actual provisions, connected by their cross-references — the AI Act here, wired to the MDR here. That web is what lets it reason about how the rules interact, not just look them up one at a time. And on this second pass, if it wrote a quote that isn't actually in the source, it catches it and fixes or flags it. So what's about to appear isn't a first guess — it's a checked answer. Here it comes."*

Beat 5 — The answer, the "aha" (2:15–2:45)

- **SCREEN:** Final answer. Slowly scroll/highlight: (1) safety-component point, (2) a

citation, (3) the GPAI note.

- **SAY:** *"And here's why the wait earns its keep. It doesn't get stuck on 'is the purpose medical' — it goes to the test that actually decides this: is the LLM a safety component of the device? That's Article 6, and the definition it turns on is right here in Article 3. And every claim links back to the provision it rests on — so this drops straight into your rationale and you verify it in seconds; you're not taking my word for it. It even flags what's easy to miss: even if it's not high-risk as a device component, an LLM is usually a general-purpose AI model, with its own obligations."*
- **The point that lands:** *"drops into your rationale, verify in seconds"* reframes the citations from a nice feature into **defensibility** — the thing an RA person actually needs.
- **Optional currency proof** (only if your answer shows the dates): *"— and notice these dates moved months ago under the Digital Omnibus, and it already has the current ones."* A quiet, concrete sign it's maintained, not a stale snapshot.

Beat 6 — Honest close + their turn (2:45–3:00)

- **SCREEN:** Rest on the citations / any uncertainty flag.
- **SAY:** *"This isn't legal advice and doesn't pretend to be — every quote is checked against the real text, and it still needs an expert to sign off. I built it on my own, for regulatory and compliance teams. So here's what I'd actually love: **the question that ate your last afternoon — send it to me, and I'll show you what CRSS does with it.** Then tell me where it's wrong."*
- **Why close on their question:** the strongest reaction to a video you *sent* isn't "nice" — it's *"let me throw my own hardest question at it."* Asking for that question is the viewer qualifying themselves, and it's the honest boundary ("expert signs off") that makes them comfortable enough to try. Don't skip the boundary to get to the ask — the boundary *is* what earns the ask.

3. Recording workflow (single take: voice + screen together)

You're recording voice and screen in one pass. It's synchronized and reads more *authentic* — exactly right for this skeptical audience (a spliced voiceover can feel like a sales reel; a real-time take feels honest). It's a little harder to land because you narrate live through the 1:40 wait, so:

1. **Capture screen + mic in one take.** macOS: **⌘+Shift+5** → **Options** → **choose your NT-1A as the microphone**, then record. (OBS or ScreenFlow work too and give you easier editing.) Speak the script as you drive.
2. **Rehearse the full run 2–3 times first.** One dry pass end-to-end — paste, narrate the

wait, reach the answer — puts the timing in muscle memory. This is what makes a single take feel calm instead of rushed.

3. **Single take ≠ no editing.** Capture in one pass, then still trim dead air, cut a fumble, or speed up a slow patch afterward. If you flub a line, **pause, breathe, and say it again** — cut the bad take later. Don't restart the whole recording over one stumble.

Riding the 1:40 wait live: the Beat 4 "how it works" narration is written to fill it — about the right length, so it should land naturally. If you run dry before the answer appears, slow down and restate the self-checking point (it bears repeating). If the answer arrives mid-sentence, finish the thought — you can nudge or trim in the edit.

NT-1A tips: quiet room; pop filter *or* sit ~15–20 cm back and speak slightly *past* it (not into it) to avoid popped "p"s; watch levels so peaks don't clip (~-12 to -6 dB). Do a 10-second record-and-listen before the real take.

Captions: most people watch muted — add them. LinkedIn can auto-generate on upload; don't let captions block you from posting. Skip the webcam for v1 (just screen + voice).

4. The optional graph glimpse (5–8s)

Worth it — it's your best visual proof that CRSS reasons over *interacting* regulations (your differentiator vs. flat "AI-for-compliance" tools). Rules: - **Frame it as regulations, not a database.** Never say "Neo4j / nodes / Cypher." Say "the actual regulations, connected by their cross-references." - **Small, legible subgraph — not a hairball.** Pre-stage a clean view: node captions set to `display_ref`, colored by label, a handful of AI Act nodes wired to MDR nodes. - **Pre-stage it; never run a live query on camera.** Have the clean graph view already open, arranged and zoomed, in another window or tab (the demo query already run, stylesheet loaded — see `neo4j_demo_visuals.md`). In your single take, just ⌘-Tab to it at the marked point, then back — you're switching to a prepared view, not building one on camera. - 5–8 seconds, then straight back. The moment it becomes about *your tech* instead of *their regulations*, you've lost them.

5. The LinkedIn post text (this is what gets it watched)

The video does the work, but the post decides whether anyone presses play. **Lead with the question, mirror their week, and ask for theirs.** Draft to adapt:

A question came up in an EU compliance community that looks simple and isn't:

If a company builds an LLM into a medical device — but for something non-medical, not diagnosis or treatment — is that LLM a "high-risk" AI system under the AI Act?

If you work in medtech RA or compliance, you know the *shape* of this one — not the exact question, but the kind that straddles the AI Act and the MDR at once, where every reading is arguable and you lose an afternoon just to be sure. Those are the afternoons RA teams lose all the time.

For the past year I've been building CRSS to turn that afternoon into a few minutes — grounded in the actual legal text, every claim traceable to its source, and it checks its own quotes before it shows you a word. Here's a ~3-minute look at it working on exactly this question.

I'm a solo developer, not a regulatory expert — so I'd value yours. **If a question like that ate your last afternoon, send it to me and I'll show you what it does** — then tell me where it gets this wrong.

[video]

Post notes: weekday morning EU time; reply to *every* comment. The two asks — "send me your question" and "tell me where it's wrong" — do double duty: they start pilot conversations *and* get you the expert validation you actually need.

6. Pre-record checklist

- [] CRSS running; Braxton's question copied, ready to paste.
- [] **Dry run done** — you've seen the real answer and it's solid (safety-component point + citations present). Never post an unverified answer to experts.
- [] Decide: let CRSS ask for the actor role (nice "it won't guess" moment) *or* add "for the manufacturer" for a straight-through answer.
- Graph view pre-staged (§4).
- [] Notifications off / Do Not Disturb on. Clean desktop, no distracting tabs.
- [] NT-1A tested; levels not clipping.
- [] Post text ready; know how you'll reply to the first few comments.

7. After you post — turn views into conversations

The recording is the top of the funnel; the *replies* are where pilots start. - **Answer every comment** thoughtfully and without defensiveness. - **When someone engages, DM them** a genuine follow-up: *"Thanks for the thoughtful comment — since you clearly live this, would you be up for putting a real question to it and telling me where it breaks?"* - Use the honest

answers below for comments/DMs:

"Isn't this just ChatGPT?" → *"A general chatbot answers from memory and invents citations. CRSS answers only from the legal text it retrieved, checks its own quotes against the source, and shows where each claim comes from — you saw it flag its own uncertainty. But tell me if it holds up under real questions."*

"How do I know it's accurate?" → *"Every claim is traceable, so you can verify it in seconds; and it checks its own quotes and flags what it can't ground. It's decision support — you still confirm. That's what a pilot is for: finding where it's wrong."*

"Is this legal advice / liability?" → *"No — it's a research and drafting aid. It doesn't replace your regulatory or legal judgment, and it's not a compliance guarantee."*

"What about our data?" → *[Answer truthfully for your setup. Safe honest version:] "It runs on public legal sources and doesn't need your confidential data to answer regulatory questions. For a pilot we'd agree exactly what, if anything, is shared."*

"Which regs does it cover?" → *"AI Act, MDR, IVDR, GDPR, the CIR 2026/977, plus AI Act and MDCG guidance. Growing."*

When you don't know: *"Honestly, that's more your world than mine — how do you handle that today?"* Turns a gap into discovery. Your strongest move, not a weakness.

8. Substance cheat-sheet (so you can speak with confidence)

Treat this as your *working understanding* — the legal text, CRSS, and a real expert are the authorities. Enough to frame the problem and understand your own demo.

The AI Act ↔ MDR connection, in plain terms

- Medical devices → **MDR (2017/745)**; IVDs → **IVDR (2017/746)**. Both are listed in the AI Act's **Annex I**.
- The AI Act has **two doors into "high-risk"**:
 - **Door 1 — Article 6(1) (product route)**: your AI is high-risk if it's a *safety component of, or is, a product covered by Annex I* (a medical device) **and** that product needs **third-party conformity assessment** (a notified body). Most Class IIa+ devices / most IVDs qualify.
 - **Door 2 — Article 6(2) + Annex III (use-case route)**: stand-alone AI in listed use cases (biometrics, employment, ...), with a limited self-assessment carve-out (Art

6(3)).

- **"Safety component" (Article 3):** a part whose *safety function*, or whose *failure*, affects health/safety. **This is the hinge of Braxton's question.**
- **Article 43:** where both apply, conformity assessment runs through the **MDR procedure** with AI-Act requirements folded in — one assessment, not two.
- **GPAI models (Chapter V):** an **LLM is usually a GPAI model**, with its *own* obligations *regardless* of the high-risk question.
- **Timelines (after the Digital Omnibus, Reg 2026/1744):** high-risk under Annex III → **2 Dec 2027**; high-risk under Annex I (device AI) → **2 Aug 2028**. (*That these just moved is itself a perfect example of the problem CRSS solves.*)

Braxton's question, explained (understand your own demo)

- **The trap:** arguing whether the LLM's *purpose* is "medical." Wrong test.
- **The real test:** is the LLM a **safety component** (Art 6(1) + Art 3)? If its function is genuinely non-safety and its failure wouldn't endanger health/safety, it's likely **not** high-risk *via the device route*.
- **Don't stop there:** check the other door (Annex III), and remember the LLM is likely a **GPAI model** with its own obligations. "Not high-risk as a device component" ≠ "no obligations."
- **The nuance that impresses experts:** it's fact-specific — how the LLM is integrated can pull it into "safety component" territory even if its purpose sounds non-medical. So the honest answer is "here's the decisive test and the provisions — an expert confirms."
- **Why it's a great demo:** the naive answers are *both* wrong; CRSS surfaces the *right question*.

Mini-glossary

- **High-risk AI system** — the AI Act tier with the heavy obligations.
- **Annex I vs III** — Annex I = existing product laws (incl. MDR/IVDR); Annex III = stand-alone high-risk use cases.
- **Conformity assessment / notified body** — the third-party check before CE marking; the AI Act plugs into it for device AI.
- **Intended purpose** — the MDR concept that fixes what a device is *for*.
- **GPAI** — general-purpose AI model (e.g. an LLM); has its own AI Act obligations.

9. Do's & don'ts

Do — open on the question (3-second rule) · mirror their recurring week · make the wait a feature · frame citations as *defensibility* · say "decision support, not legal advice" · ask for the

question that ate *their* afternoon · keep it under 3:00 · ship the rough cut.

Don't — hype or overclaim · out-expert the experts · hide the limitations · say "Neo4j/nodes/Cypher" · demo an unverified answer · let it run past 3 minutes.

10. The through-line (if you remember nothing else)

Their recurring question → **watch it work (carefully, not fast)** → **citations they could defend** → **"send me the one that ate *your* afternoon."** Make them recognise their own week, show it work credibly, be honest about the limits, hand them the mic. It's a *sent* video — its job isn't to persuade, it's to make "I need this" their own conclusion.